

1.Index:

a.

```

1
2
3 SELECT DISTINCT p.prod_id
4 FROM products p
5 JOIN inventory i
6 ON p.prod_id = i.prod_id
7 WHERE i.quan_in_stock < 100;
8
9 SELECT DISTINCT p.prod_id
10 FROM products p
11 WHERE p.prod_id IN (SELECT prod_id
12                      FROM inventory
13                      WHERE quan_in_stock < 100);
14
15 SELECT * FROM inventory WHERE quan_in_stock > 18;

```

Data Output Messages Notifications

	prod_id [PK] integer	quan_in_stock integer	sales integer
1	1	138	9
2	2	118	19
3	3	228	11
4	4	279	12
5	5	382	13
6	6	109	14
7	7	90	10
8	8	357	6
9	9	215	27
10	10	105	16

```
SELECT * FROM inventory WHERE quan_in_stock > 18;
```

b.

```
14  
15 EXPLAIN  
16 SELECT * FROM inventory WHERE quan_in_stock > 18;  
17 CREATE INDEX idx_hash ON inventory USING HASH(quan_in_stoc
```

Data Output Messages Notifications



	QUERY PLAN text	🔒
1	Seq Scan on inventory (cost=0.00..180.00 rows=9595 width=12)	
2	Filter: (quan_in_stock > 18)	

Total rows: 2 of 2

Query complete 00:00:00.044

```
CREATE INDEX idx_hash ON inventory USING  
HASH(quan_in_stock);  
EXPLAIN
```

```
SELECT * FROM inventory WHERE quan_in_stock > 18;
```

c.

```
18
19 DROP INDEX IF EXISTS idx_hash;
20
21 CREATE INDEX idx_btree ON inventory USING BTREE(quan_in_st
22 EXPLAIN
23 SELECT * FROM inventory WHERE quan_in_stock > 18;
```

Data Output Messages Notifications



	QUERY PLAN	
	text	
1	Seq Scan on inventory (cost=0.00..180.00 rows=9595 width=12)	
2	Filter: (quan_in_stock > 18)	

Total rows: 2 of 2

Query complete 00:00:00.077

```
CREATE INDEX idx_btree ON inventory USING
BTREE(quan_in_stock);
EXPLAIN
```

```
SELECT * FROM inventory WHERE quan_in_stock > 18;
```

d. There is no difference between query plans in b and c. Because numbers of products having `quan_in_stock > 18` are very large so the DBMS considered seq scan as fastest way.

e.

```
26  
27 DROP INDEX IF EXISTS idx_hash;  
28 DROP INDEX IF EXISTS idx_btree;
```

Data Output Messages Notifications

DROP INDEX

Query returned successfully in 105 msec.

Total rows: 2000 of 10000

Query complete 00:00:00.105

```
DROP INDEX IF EXISTS idx_hash;  
DROP INDEX IF EXISTS idx_btree;
```

redo b

```

14
15 EXPLAIN
16 SELECT * FROM inventory WHERE quan_in_stock > 400;
17 CREATE INDEX idx_hash ON inventory USING HASH(quan_in_stock);
18
19 CREATE INDEX idx_btree ON inventory USING BTREE(quan_in_stock);
20 EXPLAIN
21 SELECT * FROM inventory WHERE quan_in_stock > 18;
22
23 SELECT * FROM inventory
24
25 DROP INDEX IF EXISTS idx_hash;
26 DROP INDEX IF EXISTS idx_btree;

```

Data Output Messages Notifications



	QUERY PLAN	
	text	
1	Seq Scan on inventory (cost=0.00..180.00 rows=2020 width=12)	
2	Filter: (quan_in_stock > 400)	

Total rows: 2 of 2

Query complete 00:00:00.099

redo c

```

14
15 EXPLAIN
16 SELECT * FROM inventory WHERE quan_in_stock > 400;
17 CREATE INDEX idx_hash ON inventory USING HASH(quan_in_stock);
18
19 CREATE INDEX idx_btree ON inventory USING BTREE(quan_in_stock);
20 EXPLAIN
21 SELECT * FROM inventory WHERE quan_in_stock > 18;
22
23 SELECT * FROM inventory
24
25 DROP INDEX IF EXISTS idx_hash;
26 DROP INDEX IF EXISTS idx_btree;

```

Data Output Messages Notifications



	QUERY PLAN	
	text	
1	Bitmap Heap Scan on inventory (cost=27.94..108.19 rows=2020 width=12)	
2	Recheck Cond: (quan_in_stock > 400)	
3	-> Bitmap Index Scan on idx_btree (cost=0.00..27.43 rows=2020 width=0)	
4	Index Cond: (quan_in_stock > 400)	

Total rows: 4 of 4

Query complete 00:00:00.041

redo d

Now, there are differences between b and c when redo because numbers of products having `quan_in_stock > 400` are fewer.

Hash indexes only support equality comparisons when performing lookups. Bitmap index scan is useful with low cardinality and full match value.

f. There are differences between d and e when list of records in Inventory having “`quan_in_stock`” greater than 400 and 18. Because numbers of products having `quan_in_stock > 400` are fewer than numbers of products having `quan_in_stock > 18`.

2. Query:

a.

```

1
2
3 SELECT DISTINCT p.prod_id
4 FROM products p
5 JOIN inventory i
6 ON p.prod_id = i.prod_id
7 WHERE i.quan_in_stock < 100;
8
9 SELECT DISTINCT p.prod_id
10 FROM products p
11 WHERE p.prod_id IN (SELECT prod_id
12                     FROM inventory
13                     WHERE quan_in_stock < 100);

```

Data Output Messages Notifications



	prod_id [PK] integer
1	1269
2	652
3	6430
4	951
5	1898
6	70
7	5843
8	8174
9	8034
10	7662

```
SELECT DISTINCT p.prod_id  
FROM products p  
JOIN inventory i  
ON p.prod_id = i.prod_id  
WHERE i.quan_in_stock < 100;
```


```

1
2
3 SELECT DISTINCT p.prod_id
4 FROM products p
5 JOIN inventory i
6 ON p.prod_id = i.prod_id
7 WHERE i.quan_in_stock < 100;
8
9 SELECT DISTINCT p.prod_id
10 FROM products p
11 WHERE p.prod_id IN (SELECT prod_id
12                     FROM inventory
13                     WHERE quan_in_stock < 100);

```

Data Output Messages Notifications



	prod_id [PK] integer 
1	1269
2	652
3	6430
4	951
5	1898
6	70
7	5843
8	8174
9	8034
10	7662

```
SELECT DISTINCT p.prod_id
FROM products p
WHERE p.prod_id IN (SELECT prod_id
                    FROM inventory
                    WHERE quan_in_stock < 100);
```

b.

```

1
2 EXPLAIN
3 SELECT DISTINCT p.prod_id
4 FROM products p
5 JOIN inventory i
6 ON p.prod_id = i.prod_id
7 WHERE i.quan_in_stock < 100;
8
9 EXPLAIN
10 SELECT DISTINCT p.prod_id
11 FROM products p
12 WHERE p.prod_id IN (SELECT prod_id

```

Data Output Messages Notifications



	QUERY PLAN	
	text	
1	HashAggregate (cost=437.67..457.94 rows=2027 width=4)	
2	Group Key: p.prod_id	
3	-> Hash Join (cost=205.34..432.60 rows=2027 width=4)	
4	Hash Cond: (p.prod_id = i.prod_id)	
5	-> Seq Scan on products p (cost=0.00..201.00 rows=10000 width=4)	
6	-> Hash (cost=180.00..180.00 rows=2027 width=4)	
7	-> Seq Scan on inventory i (cost=0.00..180.00 rows=2027 width=4)	
8	Filter: (quan_in_stock < 100)	


```

6  ON p.prod_id = i.prod_id
7  WHERE i.quan_in_stock < 100;
8
9  EXPLAIN
10 SELECT DISTINCT p.prod_id
11 FROM products p
12 WHERE p.prod_id IN (SELECT prod_id
13                     FROM inventory
14                     WHERE quan_in_stock < 100);
15
16 EXPLAIN
17 SELECT * FROM inventory WHERE quan_in_stock > 400;

```

Data Output Messages Notifications



	QUERY PLAN	
	text	
1	HashAggregate (cost=437.67..457.94 rows=2027 width=4)	
2	Group Key: p.prod_id	
3	-> Hash Join (cost=205.34..432.60 rows=2027 width=4)	
4	Hash Cond: (p.prod_id = inventory.prod_id)	
5	-> Seq Scan on products p (cost=0.00..201.00 rows=10000 width=4)	
6	-> Hash (cost=180.00..180.00 rows=2027 width=4)	
7	-> Seq Scan on inventory (cost=0.00..180.00 rows=2027 width=4)	
8	Filter: (quan_in_stock < 100)	

No difference between 2 ways of query. Because there is no index in `quan_in_stock` (inventory), and we have index on `prod_id` of table `products`

c.

after create index `btree` on `quan_in_stock` in `inventory` table, it is used in both ways of query.

```

1
2 EXPLAIN
3 SELECT DISTINCT p.prod_id
4 FROM products p
5 JOIN inventory i
6 ON p.prod_id = i.prod_id
7 WHERE i.quan_in_stock < 100;
8
9 EXPLAIN
10 SELECT DISTINCT p.prod_id
11 FROM products p
12 WHERE p.prod_id IN (SELECT prod_id
13                     FROM inventory

```

Data Output Messages Notifications



	QUERY PLAN	
	text	
1	HashAggregate (cost=366.00..386.27 rows=2027 width=4)	
2	Group Key: p.prod_id	
3	-> Hash Join (cost=133.67..360.93 rows=2027 width=4)	
4	Hash Cond: (p.prod_id = i.prod_id)	
5	-> Seq Scan on products p (cost=0.00..201.00 rows=10000 width=4)	
6	-> Hash (cost=108.33..108.33 rows=2027 width=4)	
7	-> Bitmap Heap Scan on inventory i (cost=27.99..108.33 rows=2027 width=4)	
8	Recheck Cond: (quan_in_stock < 100)	
9	-> Bitmap Index Scan on idx_btree (cost=0.00..27.49 rows=2027 width=0)	
10	Index Cond: (quan_in_stock < 100)	


```

6  ON p.prod_id = i.prod_id
7  WHERE i.quan_in_stock < 100;
8
9  EXPLAIN
10 SELECT DISTINCT p.prod_id
11 FROM products p
12 WHERE p.prod_id IN (SELECT prod_id
13                     FROM inventory
14                     WHERE quan_in_stock < 100);
15
16 EXPLAIN
17 SELECT * FROM inventory WHERE quan_in_stock > 400;
18 CREATE INDEX idx_hash ON inventory USING HASH(quan_in_stock);

```

Data Output Messages Notifications

<div> <div>≡+</div> <div>📄</div> <div>▼</div> <div>📋</div> <div>🗑️</div> <div>🗄️</div> <div>⬇️</div> <div>📈</div> </div>	
	QUERY PLAN text 🔒
1	HashAggregate (cost=366.00..386.27 rows=2027 width=4)
2	Group Key: p.prod_id
3	-> Hash Join (cost=133.67..360.93 rows=2027 width=4)
4	Hash Cond: (p.prod_id = inventory.prod_id)
5	-> Seq Scan on products p (cost=0.00..201.00 rows=10000 width=4)
6	-> Hash (cost=108.33..108.33 rows=2027 width=4)
7	-> Bitmap Heap Scan on inventory (cost=27.99..108.33 rows=2027 width=4)
8	Recheck Cond: (quan_in_stock < 100)
9	-> Bitmap Index Scan on idx_btree (cost=0.00..27.49 rows=2027 width=0)
10	Index Cond: (quan_in_stock < 100)

Because numbers of products having $\text{quan_in_stock} < 100$ are not much. Bitmap index scan is useful with low cardinality and full match value.

3. Trigger:

```

27
28 CREATE OR REPLACE FUNCTION update_inventory()
29 RETURNS TRIGGER AS $$
30 ▼ BEGIN
31 ▼ IF TG_OP = 'INSERT' THEN
32     UPDATE inventory SET quan_in_stock = quan_in_stock - N
33     WHERE prod_id = NEW.prod_id;
34 ELSIF TG_OP = 'UPDATE' THEN
35     UPDATE inventory SET quan_in_stock = quan_in_stock + O
36     WHERE prod_id = NEW.prod_id;
37 END IF;
38 RETURN NEW;
39 END;
40 $$ LANGUAGE plpgsql;
41
42 CREATE TRIGGER update_trigger
43 AFTER INSERT OR UPDATE ON orderlines
44 FOR EACH ROW
45 EXECUTE FUNCTION update_inventory();

```

Data Output Messages Notifications

CREATE TRIGGER

Query returned successfully in 38 msec.

```

CREATE OR REPLACE FUNCTION update_inventory()
RETURNS TRIGGER AS $$
BEGIN
    IF TG_OP = 'INSERT' THEN
        UPDATE inventory SET quan_in_stock =
quan_in_stock - NEW.quantity, sales = sales +
NEW.quantity
        WHERE prod_id = NEW.prod_id;
    ELSIF TG_OP = 'UPDATE' THEN
        UPDATE inventory SET quan_in_stock =
quan_in_stock + OLD.quantity - NEW.quantity, sales
= sales + NEW.quantity - OLD.quantity
        WHERE prod_id = NEW.prod_id;
    END IF;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER update_trigger
AFTER INSERT OR UPDATE ON orderlines
FOR EACH ROW
EXECUTE FUNCTION update_inventory();

```

4. Function:

```
1 CREATE OR REPLACE FUNCTION list_product(limits INT)
2 RETURNS TABLE(prod_id INT, category INT)
3 LANGUAGE plpgsql
4 AS
5 $$
6 ▼ BEGIN
7     RETURN QUERY
8     SELECT p.prod_id, p.category
9     FROM products p
10    JOIN inventory i ON i.prod_id = p.prod_id
11   WHERE i.quan_in_stock < limits;
12 END;
13 $$
```

Data Output Messages Notifications

CREATE FUNCTION

Query returned successfully in 48 msec.


```

1 DROP FUNCTION IF EXISTS list_product;
2
3 CREATE OR REPLACE FUNCTION list_product(limits INT)
4 RETURNS TABLE(prod_id INT, category INT)
5 LANGUAGE plpgsql
6 AS
7 $$
8 ▼ BEGIN
9     RETURN QUERY
10    SELECT p.prod_id, p.category
11    FROM products p
12    JOIN inventory i ON i.prod_id = p.prod_id
13    WHERE i.quan_in_stock < limits;
14 END;
15 $$
16
17 SELECT * FROM list_product(50);|
18

```

Data Output Messages Notifications



	prod_id integer	category integer
1	12	7
2	24	3
3	39	8
4	59	7
5	74	2
6	84	3
7	117	13
8	122	9
9	135	5
10	138	10